

Ruixuan Shen  
Hoang Mai Diem Pham  
EC ENGR113DA

## MINI-PROJECT 2: Classify handwritten digits with Convolutional Neural Network (CNN)

### 1. Introduction

In the lab, we implemented the program to recognize handwritten digits using the Convolutional Neural Network model on the H7 board.

We first used python, TensorFlow, and Google Colab to train the CNN model in the cloud. The MNIST database for handwritten digits was used as the training set. The training images all have 1 channel and 28 pixels by 28 pixels for their size.

After the training, we implemented the functional modules for CNN and the CNN image classifiers on the H7 board. The CNN structure we designed is shown below:



the output width, output height, and the number of channels are shown in the screenshot below:

| Layer (type)                    | Output Shape       | Param # |
|---------------------------------|--------------------|---------|
| conv2d_2 (Conv2D)               | (None, 28, 28, 16) | 160     |
| max_pooling2d_2 (MaxPooling 2D) | (None, 14, 14, 16) | 0       |
| conv2d_3 (Conv2D)               | (None, 14, 14, 32) | 4640    |
| max_pooling2d_3 (MaxPooling 2D) | (None, 7, 7, 32)   | 0       |
| flatten_1 (Flatten)             | (None, 1568)       | 0       |
| dense_1 (Dense)                 | (None, 10)         | 15690   |
| Total params: 20,490            |                    |         |
| Trainable params: 20,490        |                    |         |
| Non-trainable params: 0         |                    |         |

For the convolution layer, we first padded the input matrix with zeros on four of the outer edges, then convolve the padded input with the filter and stride the filter across the padded matrix. After the striding, we applied ReLU to every element to eliminate the negative results.

For the pooling layer, we applied max pooling for each 2\*2 block of the input, thus returning a matrix with ½ side length on each side.

For the dense layer, after applying the neuron weight and the bias to the input, we

$$P(y = j | \mathbf{x}) = \frac{e^{\mathbf{x}^T \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^T \mathbf{w}_k}}$$

applied the softmax function for normalization and produce the discrete probability distribution vector for each digit.

The accuracy for the training in the designed structure was evaluated by python and recorded as below:

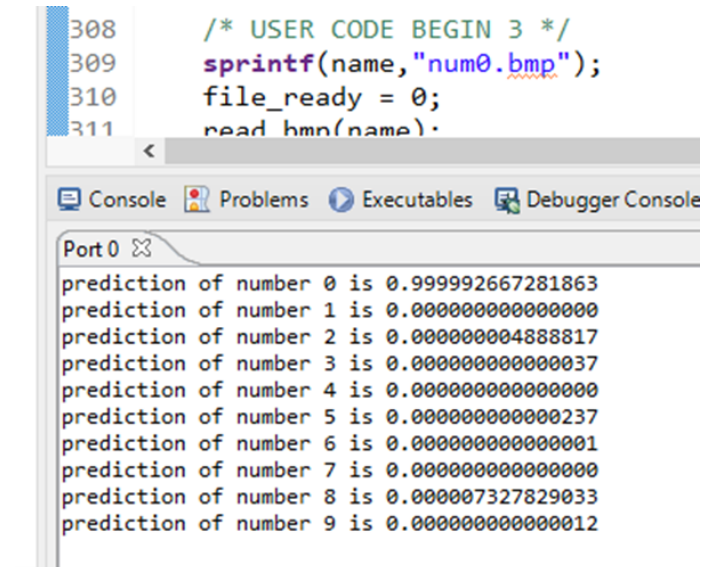
```
Accuracy of this model is:  
313/313 [=====] - 3s 10ms/step - loss: 0.1385 - accuracy: 0.9842  
[0.13846145570278168, 0.9842000007629395]
```

## 2. Result

The program failed to recognize the digit 8 from the provided test images while recognizing the rest test images successfully. The details are shown below:

### 2.1. Check number 0:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num0.bmp");
310      file_ready = 0;
311      read_bmp(name);
```



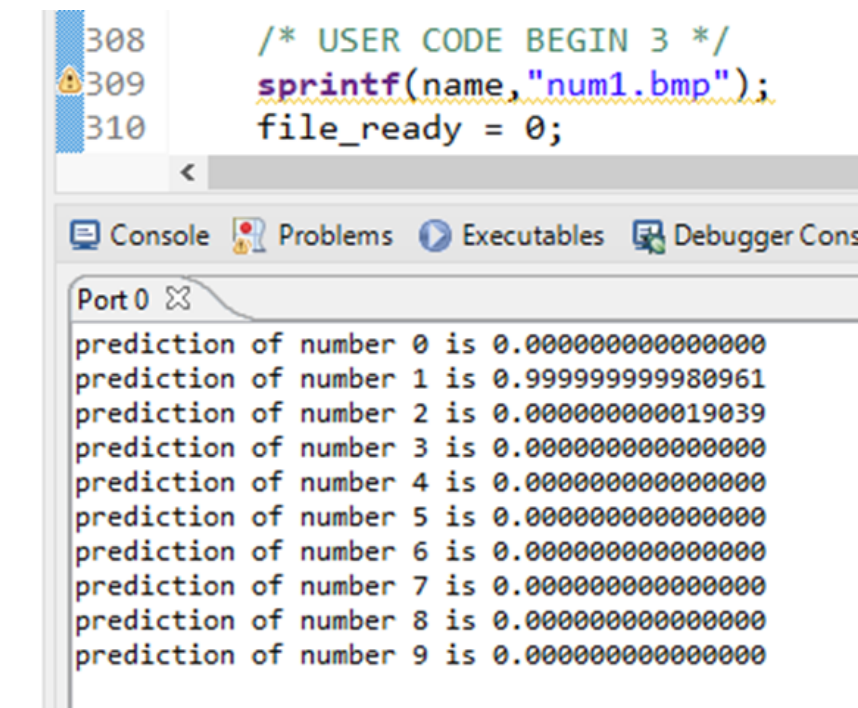
The screenshot shows an IDE with a code editor and a console window. The code editor displays lines 308 to 311 of a C program. The console window, titled 'Port 0', shows the output of the program. The output lists predictions for digits 0 through 9. The prediction for digit 0 is 0.999992667281863, which is the highest value, indicating it is the most likely digit. The prediction for digit 8 is 0.00007327829033, which is the lowest value, indicating it is the least likely digit.

Port 0

prediction of number 0 is 0.999992667281863  
prediction of number 1 is 0.000000000000000  
prediction of number 2 is 0.000000004888817  
prediction of number 3 is 0.000000000000037  
prediction of number 4 is 0.000000000000000  
prediction of number 5 is 0.000000000000237  
prediction of number 6 is 0.000000000000001  
prediction of number 7 is 0.000000000000000  
prediction of number 8 is 0.00007327829033  
prediction of number 9 is 0.000000000000012

### 2.2. Check number 1:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num1.bmp");
310      file_ready = 0;
```



The screenshot shows an IDE with a code editor and a console window. The code editor displays lines 308 to 310 of a C program. The console window, titled 'Port 0', shows the output of the program. The output lists predictions for digits 0 through 9. The prediction for digit 1 is 0.99999999980961, which is the highest value, indicating it is the most likely digit. The prediction for digit 0 is 0.000000000000000, which is the lowest value, indicating it is the least likely digit.

Port 0

prediction of number 0 is 0.000000000000000  
prediction of number 1 is 0.99999999980961  
prediction of number 2 is 0.000000000019039  
prediction of number 3 is 0.000000000000000  
prediction of number 4 is 0.000000000000000  
prediction of number 5 is 0.000000000000000  
prediction of number 6 is 0.000000000000000  
prediction of number 7 is 0.000000000000000  
prediction of number 8 is 0.000000000000000  
prediction of number 9 is 0.000000000000000

### 2.3. Check number 2:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num2.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.0000000000000000
prediction of number 2 is 1.0000000000000000
prediction of number 3 is 0.0000000000000000
prediction of number 4 is 0.0000000000000000
prediction of number 5 is 0.0000000000000000
prediction of number 6 is 0.0000000000000000
prediction of number 7 is 0.0000000000000000
prediction of number 8 is 0.0000000000000000
prediction of number 9 is 0.0000000000000000
```

### 2.4. Check number 3:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num3.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.0000000000000000
prediction of number 2 is 0.0000000000000000
prediction of number 3 is 0.999999797032286
prediction of number 4 is 0.0000000000000000
prediction of number 5 is 0.000000202967713
prediction of number 6 is 0.0000000000000000
prediction of number 7 is 0.0000000000000000
prediction of number 8 is 0.0000000000000000
prediction of number 9 is 0.0000000000000000
```

2.5. Check number 4:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num4.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.0000000000000000
prediction of number 2 is 0.0000000000000000
prediction of number 3 is 0.0000000000000000
prediction of number 4 is 1.0000000000000000
prediction of number 5 is 0.0000000000000000
prediction of number 6 is 0.0000000000000000
prediction of number 7 is 0.0000000000000000
prediction of number 8 is 0.0000000000000000
prediction of number 9 is 0.0000000000000000
```

2.6. Check number 5:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num5.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.0000000000000000
prediction of number 2 is 0.0000000000000000
prediction of number 3 is 0.0000000000000000
prediction of number 4 is 0.0000000000000000
prediction of number 5 is 1.0000000000000000
prediction of number 6 is 0.0000000000000000
prediction of number 7 is 0.0000000000000000
prediction of number 8 is 0.0000000000000000
prediction of number 9 is 0.0000000000000000
```

2.7. Check number 6:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num6.bmp");
310      file_ready = 0;
```

Port 0 X

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.0000000000000000
prediction of number 2 is 0.0000000000000000
prediction of number 3 is 0.0000000000000000
prediction of number 4 is 0.0000000000000000
prediction of number 5 is 0.000000000599480
prediction of number 6 is 0.999999999400520
prediction of number 7 is 0.0000000000000000
prediction of number 8 is 0.0000000000000000
prediction of number 9 is 0.0000000000000000
```

2.8. Check number 7:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num7.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0 X

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.000010782232172
prediction of number 2 is 0.0000000000004559
prediction of number 3 is 0.0000000000000112
prediction of number 4 is 0.0000000000000000
prediction of number 5 is 0.0000000000000000
prediction of number 6 is 0.0000000000000000
prediction of number 7 is 0.999989217763128
prediction of number 8 is 0.0000000000000029
prediction of number 9 is 0.0000000000000000
```



## 2.9. Check number 8:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num8.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.0000000000000000
prediction of number 2 is 0.0000000000000000
prediction of number 3 is 0.000000275727878
prediction of number 4 is 0.0000000000000000
prediction of number 5 is 0.000381166275216
prediction of number 6 is 0.999616711888520
prediction of number 7 is 0.0000000000000000
prediction of number 8 is 0.000001846108386
prediction of number 9 is 0.0000000000000000
```

## 2.10. Check number 9:

```
308      /* USER CODE BEGIN 3 */
309      sprintf(name, "num9.bmp");
310      file_ready = 0;
311      read_bmp(name);
```

Port 0

```
prediction of number 0 is 0.0000000000000000
prediction of number 1 is 0.000000000245994
prediction of number 2 is 0.000000001384368
prediction of number 3 is 0.009116702831457
prediction of number 4 is 0.0000000000000021
prediction of number 5 is 0.000000000031443
prediction of number 6 is 0.0000000000000000
prediction of number 7 is 0.171443969332484
prediction of number 8 is 0.000000001549626
prediction of number 9 is 0.819439324624608
```

### **3. Discussion**

We planned to have two convolution layers and two pooling layers in the first place. When we first tested our model, however, we used very small numbers for each convolution layer channel. (6 for each.) This led to many failures in the digit recognition. Only half of the test images were recognized. As we increased the number of channels, the program worked better.

Still, the program is not able to recognize 8. We tried to modify the “epochs” in the training steps as the project instruction suggested. Unfortunately, there was a limited improvement and the digit 8 was still not recognized. Further testing and modifications are needed for fixing this issue.



## 4. Code

### 4.1. main():

```
241     /* USER CODE BEGIN 3 */
242
243     sprintf(name, "num9.bmp");
244     file_ready = 0;
245     read_bmp(name);
246     if (file_ready == 1) {
247         out_img = ProcessBmp(rtext);
248         break;
249     }
250 }
251
252 //read in parameters
253 sprintf(name, "b1.txt");
254 float * b1 = read_txt(name, 16);
255
256 sprintf(name, "w1.txt");
257 float * w1 = read_txt(name, 144);
258
259 sprintf(name, "b2.txt");
260 float * b2 = read_txt(name, 32);
261
262 sprintf(name, "w2.txt");
263 float * w2 = read_txt(name, 4608);
264
265 sprintf(name, "bc.txt");
266 float * bc = read_txt(name, 10);
267
268 sprintf(name, "fc.txt");
269 float * fc = read_txt(name, 15680);
270
```

```

271 // perform NN operation
272 float *result_conv_1 = (float *)malloc(28*28*16*sizeof(float));
273 conv(out_img, w1, b1, result_conv_1, 28, 28, 1, 3, 16);
274
275 float *result_pool_1 = (float *)malloc(14*14*16*sizeof(float));
276 pool(result_conv_1, result_pool_1, 28, 28, 16, 2);
277 free(result_conv_1);
278 free(w1);
279 free(b1);
280
281 float *result_conv_2 = (float *)malloc(14*14*32*sizeof(float));
282 conv(result_pool_1, w2, b2, result_conv_2, 14, 14, 16, 3, 32);
283 free(result_pool_1);
284
285 float *result_pool_2 = (float *)malloc(7*7*32*sizeof(float));
286 pool(result_conv_2, result_pool_2, 14, 14, 32, 2);
287 free(result_conv_2);
288 free(w2);
289 free(b2);
290
291 double *result_dense = (double *)malloc(10*sizeof(double));
292 dense(result_pool_2, fc, bc, result_dense, 7*7*32, 10);
293 free(result_pool_2);
294
295 for (int i=0;i<10;i++){
296     printf("prediction of number %d is %.15f\n", i, result_dense[i]);
297 }
298 free(result_dense);
299
300 while(1);
301 /* USER CODE END 3 */
302 }

```

## 4.2. Convolution function:

```
101
102 void padding(float input[],int input_height, int input_width){
103     int input_size = input_height * input_width;
104     float input_1[input_size];
105     for (int i = 0; i < input_size; i++) {
106         input_1[i] = input[i];
107     }
108
109     int padwid = input_width + 2;
110     int padheight = input_height + 2;
111     int padsizesize = padwid * padheight;
112
113     for (int i = 0; i < padwid; i++) {
114         input[i] = 0;
115     }
116
117     for (int i = padwid * (input_height + 1); i < padsizesize; i++) {
118         input[i] = 0;
119     }
120
121     for (int i = padwid; i < padwid * (input_height + 1); i++){
122         if (i% padwid == 0 || i % padwid == padwid-1) {
123             input[i] = 0;
124         }
125         else {
126             input[i] = input_1[i-(padwid+1)-((int)(i/padwid)-1)*2];
127         }
128     }
129 }
130
```

```
131 void conv (float input[], float kernel[], float bias[], float result[], int input_height, int input_width, int input_channel, int kernel_size, int kernel_channel )
132 {
133     padding(input, input_height, input_width);
134     int result_size = input_height * input_width * kernel_channel;
135     float sum;
136     for(int n=0;n<kernel_channel;n++){
137         for (int i=0;i<input_height;i++){
138             for (int j=0;j<input_width;j++){
139                 sum=0;
140                 for(int k=0;k<kernel_channel;k++){
141                     sum+=(input[(i+0)*input_channel*kernel_channel+j*kernel_channel+k]*kernel[0*kernel_channel*kernel_channel+k*kernel_channel+n])
142                     +(input[(i+0)*16*kernel_channel+(j+1)*kernel_channel+k]*kernel[1*kernel_channel*kernel_channel+k*kernel_channel+n])
143                     +(input[(i+0)*input_channel*kernel_channel+(j+2)*kernel_channel+k]*kernel[2*kernel_channel*kernel_channel+k*kernel_channel+n]);
144                     sum+=(input[(i+1)*input_channel*kernel_channel+j*kernel_channel+k]*kernel[3*kernel_channel*kernel_channel+k*kernel_channel+n])
145                     +(input[(i+1)*16*kernel_channel+(j+1)*kernel_channel+k]*kernel[4*kernel_channel*kernel_channel+k*kernel_channel+n])
146                     +(input[(i+1)*input_channel*kernel_channel+(j+2)*kernel_channel+k]*kernel[5*kernel_channel*kernel_channel+k*kernel_channel+n]);
147                     sum+=(input[(i+2)*input_channel*kernel_channel+j*kernel_channel+k]*kernel[6*kernel_channel*kernel_channel+k*kernel_channel+n])
148                     +(input[(i+2)*16*kernel_channel+(j+1)*kernel_channel+k]*kernel[7*kernel_channel*kernel_channel+k*kernel_channel+n])
149                     +(input[(i+2)*input_channel*kernel_channel+(j+2)*kernel_channel+k]*kernel[8*kernel_channel*kernel_channel+k*kernel_channel+n]);
150                 }
151                 sum+=bias[n];
152                 result[i*input_height*kernel_channel+j*kernel_channel+n]=sum;
153             }
154         }
155     }
156
157     //ReLU: Applied to every number after convolution
158     for (int i = 0; i < result_size; i++) {
159         if (result[i] < 0) { result[i] = 0; }
160     }
161 }
162
```

### 4.3. Pooling function:

```
102
163 void pool(float input[], float output[], int height, int width, int channel, int pool_size)
164 {
165     float a, b;
166     for (int k = 0; k < (height/pool_size)*(width/pool_size)*channel ; k=k+(height/pool_size)){
167         for (int i = 0; i < (height/pool_size); i++) {
168             a = fmax(input[pool_size*(i+k)+pool_size*k], input[pool_size*(i+k)+pool_size*k+1]);
169             b = fmax(input[pool_size*(i+k)+pool_size*k + width], input[pool_size*(i+k)+pool_size*k + width + 1]);
170             output[i+k] = fmax (a,b);
171         }
172     }
173 }
```

#### 4.4. Dense function:

```
175 void dense(float input[], float kernel[], float bias[], double output[], int input_size, int output_size)
176 {
177     float y[output_size];
178     for (int i = 0; i < output_size; i++){
179         double sum = 0;
180         for (int j=0; j<input_size; j++){
181             sum = sum + kernel[i+output_size*j]*input[j];
182         }
183         y[i] = sum + bias[i];
184     }
185
186     double total = 0;
187     for (int i = 0; i < output_size; i++){
188         total = total + exp(y[i]);
189     }
190
191     for (int i = 0; i < output_size; i++){
192         output[i] = exp(y[i]) / total;
193     }
194 }
195
```

# Python Code

---


```
#import libraries for nerualnet, math and visualization
from __future__ import absolute_import, division, print_function, unicode_literals

try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras import datasets, layers, models
print(tf.__version__)

2.7.0

#set seed just for the demonstration
#tf.random.set_seed(1000);
#load in the MNIST dataset for training and testing
(train_images, train_labels), (test_images, test_labels) = datasets.mnist.load_data()

#Plot an image to see what it looks like
plt.figure()
plt.imshow(train_images[5], cmap=plt.cm.binary)
plt.grid(False)
plt.show()
```



```
#TODO: Modify the CNN structure for a slimer network
#Build the neuralnet model
#model = models.Sequential()
#model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1), padding='san
#model.add(layers.MaxPooling2D((2, 2)))
#model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
#model.add(layers.MaxPooling2D((2, 2)))
```

```

#model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
#model.add(layers.MaxPooling2D((2, 2)))
#model.add(layers.Flatten())
#model.add(layers.Dense(10, activation='softmax'))

#diem
model = models.Sequential()
model.add(layers.Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1), padding='same'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(layers.MaxPooling2D((2, 2)))
#model.add(layers.Conv2D(6, (3, 3), activation='relu', padding='same'))
#model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(10, activation='softmax'))
#Review the overall model structure
model.summary()

```

Model: "sequential\_1"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| =====                          |                    |         |
| conv2d_2 (Conv2D)              | (None, 28, 28, 16) | 160     |
| max_pooling2d_2 (MaxPooling2D) | (None, 14, 14, 16) | 0       |
| conv2d_3 (Conv2D)              | (None, 14, 14, 32) | 4640    |
| max_pooling2d_3 (MaxPooling2D) | (None, 7, 7, 32)   | 0       |
| flatten_1 (Flatten)            | (None, 1568)       | 0       |
| dense_1 (Dense)                | (None, 10)         | 15690   |
| =====                          |                    |         |
| Total params: 20,490           |                    |         |
| Trainable params: 20,490       |                    |         |
| Non-trainable params: 0        |                    |         |
| =====                          |                    |         |

```

#Review the overall model structure
#model.summary()

```

```

#Reshape the image so it can train in batch (and fit the model's input shape)
train_images = train_images.reshape((60000, 28, 28, 1))
test_images = test_images.reshape((10000, 28, 28, 1))

```

```

#Training the model
#Hint: change optimizer to 'sgd', and increase epochs if result is bad.


```



```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

```
model.fit(train_images, train_labels, epochs=20)
```


```
Epoch 1/20
1875/1875 [=====] - 42s 22ms/step - loss: 0.4031 - accuracy: 0
Epoch 2/20
1875/1875 [=====] - 43s 23ms/step - loss: 0.0764 - accuracy: 0
Epoch 3/20
1875/1875 [=====] - 44s 24ms/step - loss: 0.0581 - accuracy: 0
Epoch 4/20
1875/1875 [=====] - 44s 23ms/step - loss: 0.0494 - accuracy: 0
Epoch 5/20
1875/1875 [=====] - 43s 23ms/step - loss: 0.0414 - accuracy: 0
Epoch 6/20
1875/1875 [=====] - 44s 23ms/step - loss: 0.0405 - accuracy: 0
Epoch 7/20
1875/1875 [=====] - 44s 23ms/step - loss: 0.0383 - accuracy: 0
Epoch 8/20
1875/1875 [=====] - 44s 23ms/step - loss: 0.0301 - accuracy: 0
Epoch 9/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0295 - accuracy: 0
Epoch 10/20
1875/1875 [=====] - 44s 23ms/step - loss: 0.0295 - accuracy: 0
Epoch 11/20
1875/1875 [=====] - 44s 24ms/step - loss: 0.0278 - accuracy: 0
Epoch 12/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0258 - accuracy: 0
Epoch 13/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0274 - accuracy: 0
Epoch 14/20
1875/1875 [=====] - 44s 24ms/step - loss: 0.0234 - accuracy: 0
Epoch 15/20
1875/1875 [=====] - 44s 24ms/step - loss: 0.0214 - accuracy: 0
Epoch 16/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0228 - accuracy: 0
Epoch 17/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0228 - accuracy: 0
Epoch 18/20
1875/1875 [=====] - 45s 24ms/step - loss: 0.0198 - accuracy: 0
Epoch 19/20
1875/1875 [=====] - 44s 23ms/step - loss: 0.0248 - accuracy: 0
Epoch 20/20
1875/1875 [=====] - 43s 23ms/step - loss: 0.0219 - accuracy: 0
<keras.callbacks.History at 0x7effd1b8c4d0>
```



```
#Evaluate the performance with testing dataset
print("Accuracy of this model is:")
model.evaluate(test_images, test_labels)
```

Accuracy of this model is:

```
313/313 [=====] - 3s 10ms/step - loss: 0.1385 - accuracy: 0.98  
[0.13846145570278168, 0.9842000007629395]
```



```
#View the total number of parameters, so it doesn't overflow the LCDK's memory  
print("Total amount of parameter of model is:", model.count_params())
```

```
Total amount of parameter of model is: 20490
```

```
#Example for extract parameter form the first conv layer  
#TODO: you need to actually modify model.layers[XXXX], this XXX to fit your actually layer nu  
t1, t2 = model.layers[0].get_weights()  
np.savetxt('w1.txt', t1.flatten(), delimiter=',',fmt='%.16f')  
np.savetxt('b1.txt', t2.flatten(), delimiter=',',fmt='%.16f')
```

```
#Example for extract parameter form the second conv layer  
#TODO: you need to actually modify model.layers[XXXX], this XXX to fit your actually layer nu  
#And do it multiple times to save all the layer with parameters  
t1, t2 = model.layers[2].get_weights()  
np.savetxt('w2.txt', t1.flatten(), delimiter=',',fmt='%.16f')  
np.savetxt('b2.txt', t2.flatten(), delimiter=',',fmt='%.16f')
```

```
# t1, t2 = model.layers[4].get_weights()  
# np.savetxt('w3.txt', t1.flatten(), delimiter=',',fmt='%.16f')  
# np.savetxt('b3.txt', t2.flatten(), delimiter=',',fmt='%.16f')
```

```
t1, t2 = model.layers[5].get_weights()  
np.savetxt('fc.txt', t1.flatten(), delimiter=',',fmt='%.16f')  
np.savetxt('bc.txt', t2.flatten(), delimiter=',',fmt='%.16f')
```

```
import cv2
```


```
im = cv2.imread("num7.bmp")  
im = cv2.flip(im,0)  
im = im[:, :,0]  
for i in range (28):  
    for j in range (28):  
        im[i,j]=255 -im[i,j]
```

```
im.shape
```

```
(28, 28)
```

```
plt.imshow(im, cmap=plt.cm.binary)
```

<matplotlib.image.AxesImage at 0x7effd18bcbd0>



```
im = im.reshape(1,28,28,1)
```

```
model.predict(im)
```

```
array([[6.9705730e-35, 1.5201940e-16, 1.3742151e-22, 2.0156587e-24,  
        1.5461872e-13, 1.7622024e-19, 2.1085472e-21, 1.0000000e+00,  
        2.8086321e-19, 6.5138695e-19]], dtype=float32)
```